

Intro to Deep Learning

Adrien Aguilera--Multner

ENSEIRB-MATMECA
3A CISD

October 13, 2024

Data

- ▶ **Training Data:** Consists of input-output pairs (X, Y) .
- ▶ **Features:** The input variables X used to make predictions.
- ▶ **Labels:** The output variables Y that the model aims to predict.
- ▶ **Dataset Splitting:**
 - ▶ **Training Set:** Used to train the model.
 - ▶ **Validation Set:** Used to tune hyperparameters (optional).
 - ▶ **Test Set:** Used to evaluate the model's performance.

Multilayer Perceptron (MLP)


- ▶ **Linear layer:**

$$\mathbf{y} = \mathbf{W}\mathbf{x} + \mathbf{b} \quad (1)$$


- ▶ **Activation Functions:**

- ▶ **ReLU** (Rectified Linear Unit):

$$\text{ReLU}(\mathbf{x}) = \max(0, \mathbf{x}) \quad (2)$$



Linear activation functions produce linear decisions no matter the network size



Non-linearities allow us to approximate arbitrarily complex functions

Loss Function

- ▶ **Purpose:** Measures the distance between predicted outputs and true labels.
- ▶ **Common Loss Functions:**
 - ▶ **Mean Squared Error (MSE):**

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- ▶ **Cross-Entropy Loss:**

$$\text{Cross-Entropy} = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$$

- ▶ **Role in Training:** Guides the optimization process to minimize prediction errors.

Gradient Descent

An optimization algorithm used to minimize the loss function iteratively.

Update Rule:

$$\theta := \theta - \eta \nabla_{\theta} \mathcal{L} \quad (3)$$

where:

- ▶ θ are the model parameters.
- ▶ η is the learning rate.
- ▶ $\nabla_{\theta} \mathcal{L}$ is the gradient of the loss with respect to θ .


Variants: Adam, RMSprop, etc.

Backpropagation

How to optimize parameters ?

Process:

1. **Forward Pass:** Compute the output and loss
2. **Backward Pass:** Use chain rule to compute the gradients of the loss with respect to each parameter



3. **Parameter Update:** Apply gradient descent to minimize the loss

Types of Data

- ▶ **Images**

- ▶ Represented as pixel matrices
- ▶ 3D: [Channels, Height, Width]

- ▶ **Text**

- ▶ Represented as sequences of tokens in a vocabulary
- ▶ 1D: [Sequence Length]

- ▶ **Time Series**

- ▶ Represented as sequences over Time
- ▶ 1D: [Time]

Types of Models

- ▶ **Convolutional Neural Networks (CNNs)**
 - ▶ Mainly used for image data
 - ▶ Animation
- ▶ **Recurrent Neural Networks (RNNs)**
 - ▶ Designed for sequential data
- ▶ **Transformers**
 - ▶ Designed for handling long-range dependencies
 - ▶ Self-attention mechanism:

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{Softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_K}}\right) \mathbf{V} \quad (4)$$